



Data Mining: Concepts and Techniques

— Slides for Textbook —
— Chapter 6 & 7 —

©Jiawei Han and Micheline Kamber
Intelligent Database Systems Research Lab
School of Computing Science
Simon Fraser University, Canada
<http://www.cs.sfu.ca>



Chapter 6: Mining Association Rules in Large Databases

- Association rule mining
- Mining single-dimensional Boolean association rules from transactional databases
- Summary



Chapter 6: Mining Association Rules in Large Databases

- Association rule mining
 - Basic Concepts
 - Frequent Patterns
 - Association Rules
 - Support and Confidence
 - Road map



Chapter 6: Mining Association Rules in Large Databases

- Association rule mining
- Mining single-dimensional Boolean association rules from transactional databases
- Summary



Property of Frequent Patterns

- The **downward closure** (also called “**Apriori**”) property of frequent patterns
 - If {beer, diaper, nuts} is frequent, so is {beer, diaper}
 - Every transaction containing {beer, diaper, nuts} also contains {beer, diaper}
 - **Apriori** property: Any **subset** of a frequent itemset **must be frequent**
- Efficient mining methodology
 - If any subset of an itemset S is infrequent, then there is no chance for S to be frequent—why do we even have to consider S ? (Pruning)

Mining Association Rules—An Example

Transaction ID	Items Bought
2000	A,B,C
1000	A,C
4000	A,D
5000	B,E,F

Min. support 50%
Min. confidence 50%

Frequent Itemset	Support
{A}	75%
{B}	50%
{C}	50%
{A,C}	50%

For rule $A \Rightarrow C$:

support = support($\{A \text{ } C\}$) = 50%

confidence = support($\{A \text{ } C\}$) / support($\{A\}$) = 66.6%

The **Apriori** principle:

Any subset of a frequent itemset must be frequent



Mining Frequent Itemsets: the Key Step

- Find the *frequent itemsets*: the sets of items that have minimum support
 - A subset of a frequent itemset must also be a frequent itemset
 - i.e., if $\{AB\}$ is a frequent itemset, both $\{A\}$ and $\{B\}$ should be a frequent itemset
 - Iteratively find frequent itemsets with cardinality from 1 to k (k -itemset)
- Use the frequent itemsets to generate association rules
 - Note: If there is any itemset which is infrequent, its superset should not even be generated



The Apriori Algorithm

- **Join Step:** C_k is generated by joining L_{k-1} with itself
- **Prune Step:** Any $(k-1)$ -itemset that is not frequent cannot be a subset of a frequent k -itemset
- Pseudo-code:

C_k : Candidate itemset of size k

L_k : frequent itemset of size k

$L_1 = \{\text{frequent items}\};$

for ($k = 1; L_k \neq \emptyset; k++$) **do begin**

C_{k+1} = candidates generated from L_k ;

for each transaction t in database **do**

 increment the count of all candidates in C_{k+1}

 that are contained in t

L_{k+1} = candidates in C_{k+1} with min_support

end

return $\cup_k L_k$

The Apriori Algorithm — Example

Database D

TID	Items
100	1 3 4
200	2 3 5
300	1 2 3 5
400	2 5

Scan D

C_1

itemset	sup.
{1}	2
{2}	3
{3}	3
{4}	1
{5}	3

L_1

itemset	sup.
{1}	2
{2}	3
{3}	3
{5}	3

C_2

itemset	sup
{1 2}	1
{1 3}	2
{1 5}	1
{2 3}	2
{2 5}	3
{3 5}	2

Scan D

C_2

itemset
{1 2}
{1 3}
{1 5}
{2 3}
{2 5}
{3 5}

L_2

itemset	sup
{1 3}	2
{2 3}	2
{2 5}	3
{3 5}	2

C_3

itemset
{2 3 5}


Scan D

L_3

itemset	sup
{2 3 5}	2

Apriori: Implementation Tricks

- How to generate candidates?
 - Step 1: self-joining F_k
 - Step 2: pruning
- Example of candidate-generation
 - $F_3 = \{abc, abd, acd, ace, bcd\}$
 - Self-joining: $F_3 * F_3$
 - $abcd$ from abc and abd
 - $acde$ from acd and ace
 - Pruning:
 - $acde$ is removed because ade is not in F_3
 - $C_4 = \{abcd\}$





Mining Frequent Patterns Without Candidate Generation: FP-tree

- Compress a large database into a compact, **Frequent-Pattern tree (FP- tree)** structure
 - highly condensed, but complete for frequent pattern mining
 - avoid costly database scans
- Develop an efficient, FP-tree-based frequent pattern mining method
 - A divide-and-conquer methodology: decompose mining tasks into smaller ones
 - Avoid candidate generation: sub-database test only



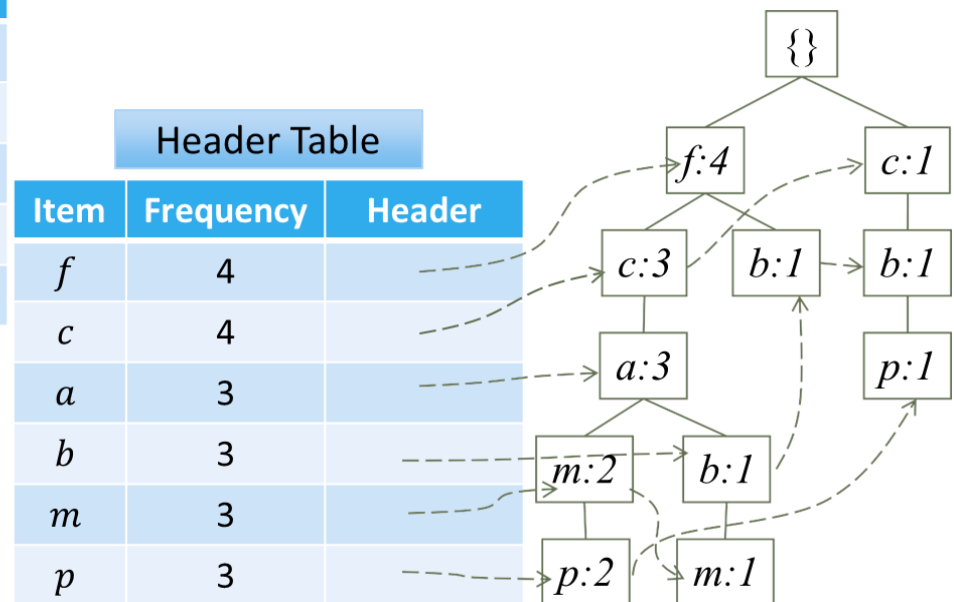
FP-Growth Method: Construction of FP-Tree

1. Create the **root of the tree**, labeled with “null”.
2. **Scan the database D** a second time. (First time scanned it to create 1-itemset and then list L).
3. The items in each transaction are processed in L order (i.e. **sorted order**).
4. A **branch is created for each transaction** with items having their support count separated by colon.
5. Whenever the same node is encountered in another transaction, just **increment the support count** of the common node or Prefix.
6. To facilitate tree traversal, an item header table is built so that each item points to its occurrences in the tree via a chain of node-links.
7. The problem of mining frequent patterns in database is transformed to that of **mining the FP-Tree**.

Example: Construct FP-tree from a Transactional DB

TID	Items in the Transaction	Ordered, frequent items
100	{f, a, c, d, g, i, m, p}	{f, c, a, m, p}
200	{a, b, c, f, l, m, o}	{f, c, a, b, m}
300	{b, f, h, j, o, w}	{f, b}
400	{b, c, k, s, p}	{c, b, p}
500	{a, f, c, e, l, p, m, n}	{f, c, a, m, p}


1. Scan DB once, find single item frequent pattern: Let min_sup = 3
f:4, a:3, c:4, b:3, m:3, p:3
2. Sort frequent items in frequency descending order, f-list
F-list = f-c-a-b-m-p
3. Scan DB again, construct FP-tree




Constructing the FP-Tree

TID	Items in the Transaction	Ordered, frequent items
100	{f, a, c, d, g, i, m, p}	{f, c, a, m, p}
200	{a, b, c, f, l, m, o}	{f, c, a, b, m}
300	{b, f, h, j, o, w}	{f, b}
400	{b, c, k, s, p}	{c, b, p}
500	{a, f, c, e, l, p, m, n}	{f, c, a, m, p}

Item	Frequency	Header
f	4	
c	4	
a	3	
b	3	
m	3	
p	3	



{f, c, a, m, p}




{f, c, a, b, m}


Constructing the FP-Tree

TID	Items in the Transaction	Ordered, frequent items
100	{f, a, c, d, g, i, m, p}	{f, c, a, m, p}
200	{a, b, c, f, l, m, o}	{f, c, a, b, m}
300	{b, f, h, j, o, w}	{f, b}
400	{b, c, k, s, p}	{c, b, p}
500	{a, f, c, e, l, p, m, n}	{f, c, a, m, p}

Item	Frequency	Header
f	4	
c	4	
a	3	
b	3	
m	3	
p	3	



{f, b}




{c, b, p}

Constructing the FP-Tree

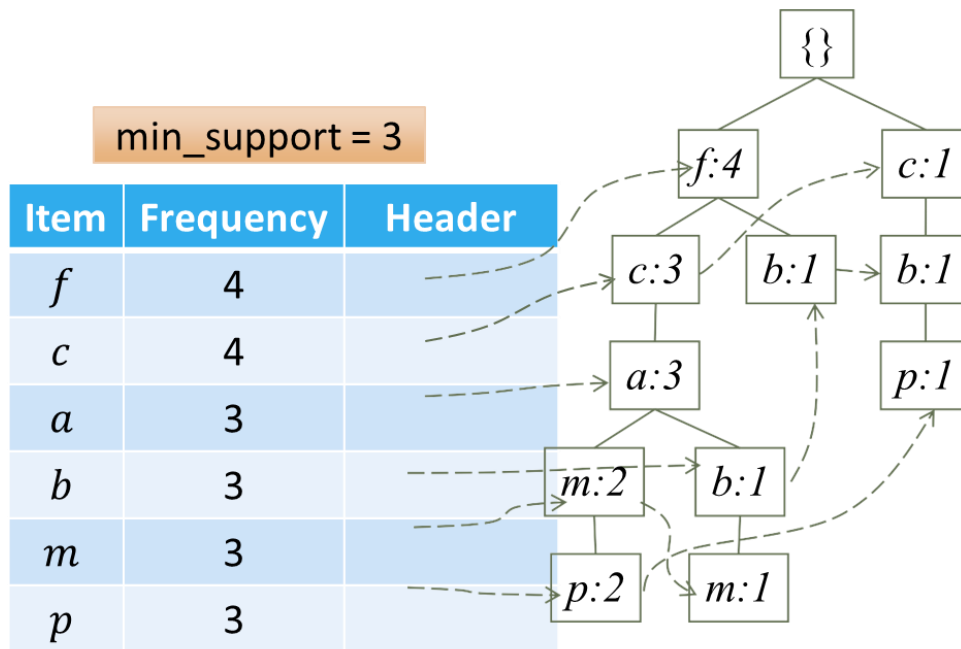
TID	Items in the Transaction	Ordered, frequent items
100	{f, a, c, d, g, i, m, p}	{f, c, a, m, p}
200	{a, b, c, f, l, m, o}	{f, c, a, b, m}
300	{b, f, h, j, o, w}	{f, b}
400	{b, c, k, s, p}	{c, b, p}
500	{a, f, c, e, l, p, m, n}	{f, c, a, m, p}

Item	Frequency	Header
f	4	
c	4	
a	3	
b	3	
m	3	
p	3	



Generate Conditional Pattern Bases

- Pattern mining can be partitioned according to current patterns
 Patterns containing p: p's conditional database: fcam:2, cb:1
 Patterns having m but no p : m's conditional database: fca:2, fcab:1




Conditional pattern bases

Item	Conditional pattern base
<i>c</i>	<i>f</i> : 3
<i>a</i>	<i>fc</i> : 3
<i>b</i>	<i>fca</i> : 1, <i>f</i> : 1, <i>c</i> : 1
<i>m</i>	<i>fca</i> : 2, <i>fcab</i> : 1
<i>p</i>	<i>fcam</i> : 2, <i>cb</i> : 1

Generate Conditional FP-Tree

- Calculate recurring items in conditional pattern base following each item




min_support = 3

Item	Conditional pattern base	Conditional FP-Tree
<i>c</i>	<i>f</i> :3	<i>f</i> :3
<i>a</i>	<i>f</i> <i>c</i> :3	<i>f</i> :3, <i>c</i> :3
<i>b</i>	<i>f</i> <i>c</i> <i>a</i> :1, <i>f</i> :1, <i>c</i> :1	<i>f</i> :2, <i>c</i> :2, <i>a</i> :1
<i>m</i>	<i>f</i> <i>c</i> <i>a</i> :2, <i>f</i> <i>c</i> <i>a</i> <i>b</i> :1	<i>f</i> :3, <i>c</i> :3, <i>a</i> :3, <i>b</i> :1
<i>p</i>	<i>f</i> <i>c</i> <i>a</i> <i>m</i> :2, <i>c</i> <i>b</i> :1	<i>f</i> :2, <i>c</i> :3, <i>a</i> :2, <i>m</i> :2, <i>b</i> :1

Generate Frequent Patterns

- Get itemsets for each conditional FP-tree




min_support = 3

Item	Conditional pattern base	Conditional FP-Tree	Frequent Patterns
c	f:3	f:3	{f,c:3},{c:3}
a	fc:3	f:3, c:3	{f,c,a:3},{f,a:3},{c,a:3},{a:3}
b	fca:1, f:1, c:1	f:2, c:2, a:1	{b:3}
m	fca:2, fcab:1	f:3, c:3, a:3, b:1	{f,c,a,m:3},{f,c,m:3}, {f,a,m:3},{c,a,m:3},{f,m:3}, {c,m:3},{a,m:3},{m:3}
p	fcam:2, cb:1	f:2, c:3, a:2, m:2, b:1	{c,p:3},{p:3}

Generate Association Rules

- Create rules for each frequent itemset



Frequent Patterns

$\{f, c:3\}, \{c:3\}$

$\{f, c, a:3\}, \{f, a:3\}, \{c, a:3\}, \{a:3\}$

$\{b:3\}$

$\{f, c, a, m:3\}, \{f, c, m:3\},$
 $\{f, a, m:3\}, \{c, a, m:3\}, \{f, m:3\},$
 $\{c, m:3\}, \{a, m:3\}, \{m:3\}$

$\{c, p:3\}, \{p:3\}$

Consider this itemset: $\{f, c, a:3\}$

Generate all subsets and rules:

$f \rightarrow c \wedge a$

$c \rightarrow f \wedge a$

$a \rightarrow f \wedge c$

$f \wedge c \rightarrow a$

$f \wedge a \rightarrow c$

$c \wedge a \rightarrow f$

~~$f \wedge c \wedge a$~~

Calculate Confidence

- Calculate confidence for each rule and check minimum

Item	Frequency
<i>f</i>	4
<i>c</i>	4
<i>a</i>	3
<i>b</i>	3
<i>m</i>	3
<i>p</i>	3

Frequent Patterns

$\{f, c:3\}$

$\{f, c, a:3\}, \{f, a:3\}, \{c, a:3\}, \{a:3\}$

$X \rightarrow Y$

confidence, *c*, **conditional probability** that a transaction having *X* also contains *Y*

$$C = (\text{count of } X \& Y) / (\text{Count of } X)$$

Consider this itemset: $\{f, c, a:3\}$

Generate all subsets and rules:

$f \rightarrow c^{\wedge}a$ $\text{conf.} = 3/4$

$c \rightarrow f^{\wedge}a$ $\text{conf.} = 3/4$

$a \rightarrow f^{\wedge}c$ $\text{conf.} = 3/3$

$f^{\wedge}c \rightarrow a$ $\text{conf.} = 3/3$

$f^{\wedge}a \rightarrow c$ $\text{conf.} = 3/3$

$c^{\wedge}a \rightarrow f$ $\text{conf.} = 3/3$

~~$f^{\wedge}c^{\wedge}a$~~



How to Judge if a Rule/Pattern Is Interesting?

- Pattern-mining will generate a large set of patterns/rules
 - Not all the generated patterns/rules are interesting
- **Interestingness measures: Objective vs. subjective**
 - **Objective** interestingness measures: Based on threshold values controlled by the user.
 - Support, confidence, correlation, ...
 - **Subjective** interestingness measures: Often based on earlier user experiences and beliefs
 - Query-based: Relevant to a user's particular request
 - Against one's knowledge-base: unexpected, freshness, timeliness
 - Visualization tools: Multi-dimensional, interactive examination



Support and confidence

- If confidence gets a value of 100 % the rule is an **exact rule**
- Even if confidence reaches high values the rule is not useful unless the support value is high as well
- Rules that have both high confidence and support are called **strong rules**
- But strong rules are not necessarily interesting.

Limitation of the Support-Confidence Framework

- Are s and c interesting in association rules: " $A \Rightarrow B$ " [s,c]?
- Example: Suppose one school may have the following statistics on # of students who may play basketball and/or eat cereal:

	play-basketball	not play-basketball	sum (row)
eat-cereal	400	350	750
not eat-cereal	200	50	250
sum (col.)	600	400	1000

2-way contingency table

- Association rule mining may generate the following:
- $\text{play-basketball} \Rightarrow \text{eat-cereal}$ [40%, 66.7%] (higher s & c)
- Looks good. But if you generate another rule
- $\neg \text{play-basketball} \Rightarrow \text{eat-cereal}$ [35%, 87.5%] (high s & c)
- These two rules confuse the cereal company.



Interestingness Measure: Lift

- Measure of dependent/correlated events: lift
 $\text{lift}(B,C) = (c(B \rightarrow C)) / (s(C)) = (s(B \cup C)) / (s(B) \times s(C))$
- $\text{lift}(B,C)$ may tell how B and C are correlated
 - $\text{lift}(B,C) = 1$: B and C are independent
 - > 1 : positively correlated
 - < 1 : negatively correlated

Lift is more telling than s & c

	B	$\neg B$	Σ_{row}
C	400	350	750
$\neg C$	200	50	250
Σ_{col}	600	400	1000

- For our example,
 $\text{lift}(B,C) = (400/1000) / (600/1000 \times 750/1000) = 0.89$
 $\text{lift}(B,\neg C) = (200/1000) / (600/1000 \times 250/1000) = 1.33$
- Thus, B and C are negatively correlated since $\text{lift}(B, C) < 1$;
 - B and $\neg C$ are positively correlated since $\text{lift}(B, \neg C) > 1$

Is Lift Always A Good Measure?

- Null transactions: Transactions that contain neither B nor C
- Let's examine the dataset D
- BC (100) is much rarer than $B\neg C$ (1000) and $\neg BC$ (1000), but there are many $\neg B\neg C$ (100000)
- Unlikely B & C will happen together!
- But, $\text{Lift}(B, C) = 8.44 \gg 1$ (Lift shows B and C are strongly positively correlated!)

	B	$\neg B$	Σ_{row}
C	100	1000	1100
$\neg C$	1000	100000	101000
Σ_{col}	1100	101000	102100

null transactions

Interestingness Measures & Null-Invariance

- Null invariance: Value does not change with the # of null-transactions
- A few interestingness measures: Some are null invariant

Measure	Definition	Range	Null-Invariant
$\chi^2(A, B)$	$\sum_{i,j=0,1} \frac{(e(a_i b_j) - o(a_i b_j))^2}{e(a_i b_j)}$	$[0, \infty]$	No
$Lift(A, B)$	$\frac{s(A \cup B)}{s(A) \times s(B)}$	$[0, \infty]$	No
$AllConf(A, B)$	$\frac{s(A \cup B)}{\max\{s(A), s(B)\}}$	$[0, 1]$	Yes
$Jaccard(A, B)$	$\frac{s(A \cup B)}{s(A) + s(B) - s(A \cup B)}$	$[0, 1]$	Yes
$Cosine(A, B)$	$\frac{s(A \cup B)}{\sqrt{s(A) \times s(B)}}$	$[0, 1]$	Yes
$Kulczynski(A, B)$	$\frac{1}{2} \left(\frac{s(A \cup B)}{s(A)} + \frac{s(A \cup B)}{s(B)} \right)$	$[0, 1]$	Yes
$MaxConf(A, B)$	$\max\left\{ \frac{s(A)}{s(A \cup B)}, \frac{s(B)}{s(A \cup B)} \right\}$	$[0, 1]$	Yes

χ^2 and lift are not null-invariant

Jaccard, cosine, AllConf, MaxConf, and Kulczynski are null-invariant measures

ExKul: 0- negatively correlated, 0.5- neutral, 1- positively correlated

Null Invariance: An Important Property

milk vs. coffee contingency table

	milk	\neg milk	Σ_{row}
coffee	mc	$\neg mc$	c
\neg coffee	$m\neg c$	$\neg m\neg c$	$\neg c$
Σ_{col}	m	$\neg m$	Σ

Dataset	mc	$\neg mc$	$m\neg c$	$\neg m\neg c$
D_1	10,000	1,000	1,000	100,000
D_2	10,000	1,000	1,000	100
D_3	100	1,000	1,000	100,000
D_4	1,000	1,000	1,000	100,000
D_5	1,000	100	10,000	100,000
D_6	1,000	10	100,000	100,000

- Let's look at another ex. Check the first 4 data sets.
- m and c are **positively associated** in D_1 and D_2 , because $mc(10,000)$ is considerably greater than $m\bar{c}(1000)$ and $\bar{m}c(1000)$
- **Negatively associated** in D_3 , because $mc(100)$ is considerably lesser than $m\bar{c}(1000)$ and $\bar{m}c(1000)$
- **Neutral** in D_4 , because $mc(1000)$ is equal to $m\bar{c}(1000)$ and $\bar{m}c(1000)$

Null Invariance: An Important Property

- Why is null invariance crucial for the analysis of massive transaction data?
 - Many transactions may contain neither milk nor coffee

milk vs. coffee contingency table

	milk	\neg milk	Σ_{row}
coffee	mc	$\neg mc$	c
\neg coffee	$m\neg c$	$\neg m\neg c$	$\neg c$
Σ_{col}	m	$\neg m$	Σ

- Lift is not null-invariant: not good to evaluate data that contain too many (D1) or too few (D2) null transactions
- Many measures are not null-invariant

Dataset	mc	$\neg mc$	$m\neg c$	$\neg m\neg c$	χ^2	$Lift$
D_1	10,000	1,000	1,000	100,000	90557	9.26
D_2	10,000	1,000	1,000	100	0	1
D_3	100	1,000	1,000	100,000	670	8.44
D_4	1,000	1,000	1,000	100,000	24740	25.75
D_5	1,000	100	10,000	100,000	8173	9.18
D_6	1,000	10	100,000	100,000	965	1.97

Null-transactions
w.r.t. m and c

Comparison of Null-Invariant Measures

- Not all null-invariant measures are created equal
- D4-D6 differentiate the null-invariant measures
- Imbalance Ratio (IR) can measure which is better

2-variable contingency table

	milk	\neg milk	Σ_{row}
coffee	mc	$\neg mc$	c
\neg coffee	$m\neg c$	$\neg m\neg c$	$\neg c$
Σ_{col}	m	$\neg m$	Σ

All 5 are null-invariant

Dataset	mc	$\neg mc$	$m\neg c$	$\neg m\neg c$	<i>AllConf</i>	<i>Jaccard</i>	<i>Cosine</i>	<i>Kulc</i>	<i>MaxConf</i>
D_1	10,000	1,000	1,000	100,000	0.91	0.83	0.91	0.91	0.91
D_2	10,000	1,000	1,000	100	0.91	0.83	0.91	0.91	0.91
D_3	100	1,000	1,000	100,000	0.09	0.05	0.09	0.09	0.09
D_4	1,000	1,000	1,000	100,000	0.5	0.33	0.5	0.5	0.5
D_5	1,000	100	10,000	100,000	0.09	0.09	0.29	0.5	0.91
D_6	1,000	10	100,000	100,000	0.01	0.01	0.10	0.5	0.99

Subtle: They disagree on most cases



What Measures to Choose for Effective Pattern Evaluation?

- Null value cases are predominant in many large datasets
- Neither milk nor coffee is in most of the baskets; neither Mike nor Jim is an author in most of the papers;
- Null-invariance is an important property
- Lift, χ^2 and cosine are good measures if null transactions are not predominant
- Otherwise, choose others to judge the interestingness of a pattern (e.g. Kulczynski + Imbalance Ratio)



Summary

- Basic Concepts:
 - Frequent Patterns, Association Rules, Closed Patterns and Max-Patterns
- Frequent Itemset Mining Methods
 - The Downward Closure Property and The Apriori Algorithm
 - FPGrowth: A Frequent Pattern-Growth Approach
- Which Patterns Are Interesting?—Pattern Evaluation Methods
 - Interestingness Measures: Lift and χ^2
 - Null-Invariant Measures
 - Comparison of Interestingness Measures



References

- R. Agarwal, C. Aggarwal, and V. V. V. Prasad. A tree projection algorithm for generation of frequent itemsets. In Journal of Parallel and Distributed Computing (Special Issue on High Performance Data Mining), 2000.
- R. Agrawal, T. Imielinski, and A. Swami. Mining association rules between sets of items in large databases. SIGMOD'93, 207-216, Washington, D.C.
- R. Agrawal and R. Srikant. Fast algorithms for mining association rules. VLDB'94 487-499, Santiago, Chile.
- R. Agrawal and R. Srikant. Mining sequential patterns. ICDE'95, 3-14, Taipei, Taiwan.
- R. J. Bayardo. Efficiently mining long patterns from databases. SIGMOD'98, 85-93, Seattle, Washington.
- S. Brin, R. Motwani, and C. Silverstein. Beyond market basket: Generalizing association rules to correlations. SIGMOD'97, 265-276, Tucson, Arizona.
- S. Brin, R. Motwani, J. D. Ullman, and S. Tsur. Dynamic itemset counting and implication rules for market basket analysis. SIGMOD'97, 255-264, Tucson, Arizona, May 1997.
- K. Beyer and R. Ramakrishnan. Bottom-up computation of sparse and iceberg cubes. SIGMOD'99, 359-370, Philadelphia, PA, June 1999.
- D.W. Cheung, J. Han, V. Ng, and C.Y. Wong. Maintenance of discovered association rules in large databases: An incremental updating technique. ICDE'96, 106-114, New Orleans, LA.
- M. Fang, N. Shivakumar, H. Garcia-Molina, R. Motwani, and J. D. Ullman. Computing iceberg queries efficiently. VLDB'98, 299-310, New York, NY, Aug. 1998.



References (2)

- G. Grahne, L. Lakshmanan, and X. Wang. Efficient mining of constrained correlated sets. ICDE'00, 512-521, San Diego, CA, Feb. 2000.
- Y. Fu and J. Han. Meta-rule-guided mining of association rules in relational databases. KDOOD'95, 39-46, Singapore, Dec. 1995.
- T. Fukuda, Y. Morimoto, S. Morishita, and T. Tokuyama. Data mining using two-dimensional optimized association rules: Scheme, algorithms, and visualization. SIGMOD'96, 13-23, Montreal, Canada.
- E.-H. Han, G. Karypis, and V. Kumar. Scalable parallel data mining for association rules. SIGMOD'97, 277-288, Tucson, Arizona.
- J. Han, G. Dong, and Y. Yin. Efficient mining of partial periodic patterns in time series database. ICDE'99, Sydney, Australia.
- J. Han and Y. Fu. Discovery of multiple-level association rules from large databases. VLDB'95, 420-431, Zurich, Switzerland.
- J. Han, J. Pei, and Y. Yin. Mining frequent patterns without candidate generation. SIGMOD'00, 1-12, Dallas, TX, May 2000.
- T. Imielinski and H. Mannila. A database perspective on knowledge discovery. Communications of ACM, 39:58-64, 1996.
- M. Kamber, J. Han, and J. Y. Chiang. Metarule-guided mining of multi-dimensional association rules using data cubes. KDD'97, 207-210, Newport Beach, California.
- M. Klemettinen, H. Mannila, P. Ronkainen, H. Toivonen, and A.I. Verkamo. Finding interesting rules from large sets of discovered association rules. CIKM'94, 401-408, Gaithersburg, Maryland.



References (3)

- F. Korn, A. Labrinidis, Y. Kotidis, and C. Faloutsos. Ratio rules: A new paradigm for fast, quantifiable data mining. VLDB'98, 582-593, New York, NY.
- B. Lent, A. Swami, and J. Widom. Clustering association rules. ICDE'97, 220-231, Birmingham, England.
- H. Lu, J. Han, and L. Feng. Stock movement and n-dimensional inter-transaction association rules. SIGMOD Workshop on Research Issues on Data Mining and Knowledge Discovery (DMKD'98), 12:1-12:7, Seattle, Washington.
- H. Mannila, H. Toivonen, and A. I. Verkamo. Efficient algorithms for discovering association rules. KDD'94, 181-192, Seattle, WA, July 1994.
- H. Mannila, H Toivonen, and A. I. Verkamo. Discovery of frequent episodes in event sequences. Data Mining and Knowledge Discovery, 1:259-289, 1997.
- R. Meo, G. Psaila, and S. Ceri. A new SQL-like operator for mining association rules. VLDB'96, 122-133, Bombay, India.
- R.J. Miller and Y. Yang. Association rules over interval data. SIGMOD'97, 452-461, Tucson, Arizona.
- R. Ng, L. V. S. Lakshmanan, J. Han, and A. Pang. Exploratory mining and pruning optimizations of constrained associations rules. SIGMOD'98, 13-24, Seattle, Washington.
- N. Pasquier, Y. Bastide, R. Taouil, and L. Lakhal. Discovering frequent closed itemsets for association rules. ICDT'99, 398-416, Jerusalem, Israel, Jan. 1999.



References (4)

- J.S. Park, M.S. Chen, and P.S. Yu. An effective hash-based algorithm for mining association rules. SIGMOD'95, 175-186, San Jose, CA, May 1995.
- J. Pei, J. Han, and R. Mao. CLOSET: An Efficient Algorithm for Mining Frequent Closed Itemsets. DMKD'00, Dallas, TX, 11-20, May 2000.
- J. Pei and J. Han. Can We Push More Constraints into Frequent Pattern Mining? KDD'00. Boston, MA. Aug. 2000.
- G. Piatetsky-Shapiro. Discovery, analysis, and presentation of strong rules. In G. Piatetsky-Shapiro and W. J. Frawley, editors, Knowledge Discovery in Databases, 229-238. AAAI/MIT Press, 1991.
- B. Ozden, S. Ramaswamy, and A. Silberschatz. Cyclic association rules. ICDE'98, 412-421, Orlando, FL.
- J.S. Park, M.S. Chen, and P.S. Yu. An effective hash-based algorithm for mining association rules. SIGMOD'95, 175-186, San Jose, CA.
- S. Ramaswamy, S. Mahajan, and A. Silberschatz. On the discovery of interesting patterns in association rules. VLDB'98, 368-379, New York, NY.
- S. Sarawagi, S. Thomas, and R. Agrawal. Integrating association rule mining with relational database systems: Alternatives and implications. SIGMOD'98, 343-354, Seattle, WA.
- A. Savasere, E. Omiecinski, and S. Navathe. An efficient algorithm for mining association rules in large databases. VLDB'95, 432-443, Zurich, Switzerland.
- A. Savasere, E. Omiecinski, and S. Navathe. Mining for strong negative associations in a large database of customer transactions. ICDE'98, 494-502, Orlando, FL, Feb. 1998.



References (5)

- C. Silverstein, S. Brin, R. Motwani, and J. Ullman. Scalable techniques for mining causal structures. VLDB'98, 594-605, New York, NY.
- R. Srikant and R. Agrawal. Mining generalized association rules. VLDB'95, 407-419, Zurich, Switzerland, Sept. 1995.
- R. Srikant and R. Agrawal. Mining quantitative association rules in large relational tables. SIGMOD'96, 1-12, Montreal, Canada.
- R. Srikant, Q. Vu, and R. Agrawal. Mining association rules with item constraints. KDD'97, 67-73, Newport Beach, California.
- H. Toivonen. Sampling large databases for association rules. VLDB'96, 134-145, Bombay, India, Sept. 1996.
- D. Tsur, J. D. Ullman, S. Abitboul, C. Clifton, R. Motwani, and S. Nestorov. Query flocks: A generalization of association-rule mining. SIGMOD'98, 1-12, Seattle, Washington.
- K. Yoda, T. Fukuda, Y. Morimoto, S. Morishita, and T. Tokuyama. Computing optimized rectilinear regions for association rules. KDD'97, 96-103, Newport Beach, CA, Aug. 1997.
- M. J. Zaki, S. Parthasarathy, M. Ogihara, and W. Li. Parallel algorithm for discovery of association rules. Data Mining and Knowledge Discovery, 1:343-374, 1997.
- M. Zaki. Generating Non-Redundant Association Rules. KDD'00. Boston, MA. Aug. 2000.
- O. R. Zaiane, J. Han, and H. Zhu. Mining Recurrent Items in Multimedia with Progressive Resolution Refinement. ICDE'00, 461-470, San Diego, CA, Feb. 2000.

<http://www.cs.sfu.ca/~han/dmbook>



Thank you !!!